


SOLO cuts the time before a one-token decision

1,088 ContractNLI decisions · JEV-9B BF16 · RTX 4090 · concurrency 4 · prefix cache on



Median prefill / first-token interval: 726.1 → 219.3 ms (3.31×)

Post-first-token decode: 0.0 ms in all 2,176 measurements · 1 output token per request